

Geometric Data Analysis HW 1

Gilad Turok, gt2453
gt2453@columbia.edu

February 21, 2023

1 *k*-means vs Single-Linkage Clustering

I generated data with three 2-dimensional Gaussians with identity covariance matrices. I cluster this data with *k*-means and single-linkage clustering with means that vary. I use 4 different sets of means, beginning from very close to each other to further and further apart:

$$\begin{aligned}\mu_1 &= \begin{bmatrix} 0 \\ 0 \end{bmatrix}, & \mu_2 &= \begin{bmatrix} 0 \\ 0 \end{bmatrix}, & \mu_3 &= \begin{bmatrix} 0 \\ 0 \end{bmatrix} \\ \mu_1 &= \begin{bmatrix} -2 \\ 2 \end{bmatrix}, & \mu_2 &= \begin{bmatrix} 0 \\ -2 \end{bmatrix}, & \mu_3 &= \begin{bmatrix} 2 \\ 2 \end{bmatrix} \\ \mu_1 &= \begin{bmatrix} -5 \\ 5 \end{bmatrix}, & \mu_2 &= \begin{bmatrix} 0 \\ -5 \end{bmatrix}, & \mu_3 &= \begin{bmatrix} 5 \\ 5 \end{bmatrix} \\ \mu_1 &= \begin{bmatrix} -10 \\ 10 \end{bmatrix}, & \mu_2 &= \begin{bmatrix} 10 \\ -10 \end{bmatrix}, & \mu_3 &= \begin{bmatrix} 10 \\ 10 \end{bmatrix}\end{aligned}$$

The *k*-means algorithm is very sensitive to initialization of cluster centers while single-linkage clustering is not – it converges to the same result. When the means are all the same, i.e. $\mu_1 = \mu_2 = \mu_3 = 0$, the true clusters are all on top of each other, as shown in the top left plot. *k*-means clusters this into 3 clusters while single-linkage clustering places nearly all data points in one cluster. In general, when the means are close together, single-linkage clustering often collapses into one entire cluster while *k*-means does not.

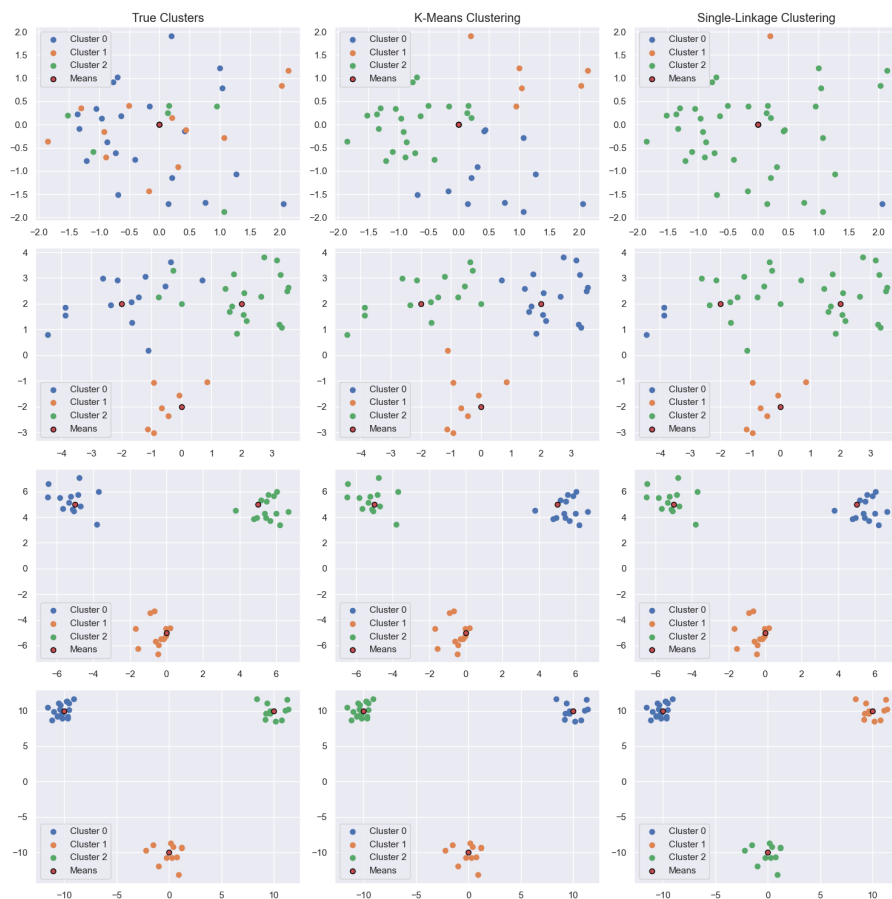


Figure 1: k -means vs single-linkage clustering on 3 Gaussians with means at different distances

2 k -means Clustering With Noise

2.1 Gaussian Noise

I added Gaussian noise $\mathcal{N}(0, \sigma^2)$ to my data \mathbf{X} . I plot the true clusters and set σ^2 to the following values: 0, 0.5, 1, 2, 5. When $\sigma^2 = 0$, k -means performs well, correctly clustering nearly all the data points. However, as the noise increases, k -means becomes less and less accurate. Due to the random initialization of the means, k -means is very sensitive to noise – this is why multiple runs of k -means is often recommended.

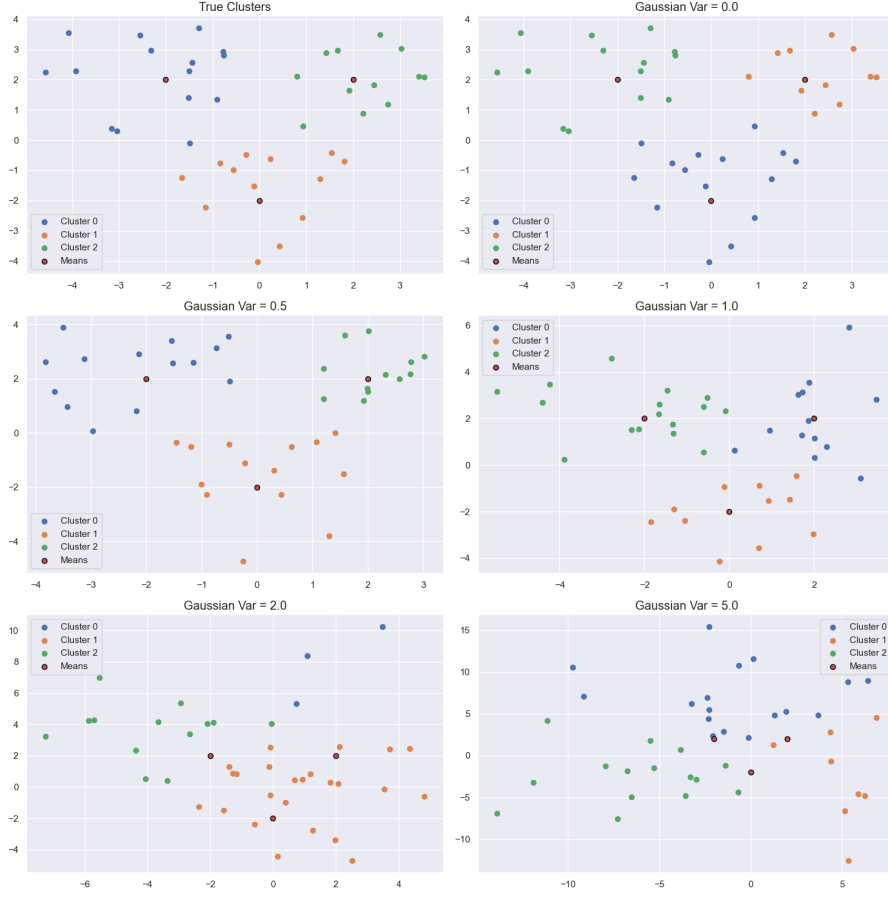


Figure 2: k -means with increasing amounts of Gaussian noise

2.2 Adversarial Noise

Adversarial noise in single-linkage clustering can be used to create a "bridge" between two clusters that are otherwise far apart. In k -means clustering, for a given cluster-center initialization, adversarial points may also sometimes create a "bridge" that slowly shifts the cluster center towards the adversarial point as the algorithm converges. However, this problem is much severe than in the case of single-linkage clustering. Furthermore, adversarial noise can be used to place all L data points into their own cluster which k -means would miss. Lastly, adversarial noise can challenge k -means clustering by creating a uniform grid of points, if $|L|$ is sufficiently large enough; this would render the entire clustering procedure useless and unusable.

3 Hierarchial k -Means and k -Medians

One can make a hierarchial version of k -means or k -medians as k varies. However, there are many challenges that arise from the non-deterministic initialization of these algorithms. If there was a fixed rule, this would alleviate such problems.

Nonetheless, there are interesting extensions that can create a form of hierarchial k -means or k -median algorithms – I will elaborate upon one such way. Let there be k desired clusters and n data points $x_1 \dots x_n$.

Hierarchial clustering e.g. single-linkage clustering suffers from expensive computational overhead. If one knows the general range of desired clusters, i.e. $k = [4, 10]$, hierarchial clustering begins from $k = n$, where n can be tens or hundreds of thousands of data points. Instead, one can run k -means or k -medians with $k = 10$ and merge the nearest clusters as per single-linkage clustering. This would allow one to more efficeintly create a dendrogram with only $k = [4, 10]$, for example. However, to ensure robustness, one would want to initialize k -means and k -medians multiple times and compare the results.

4 Clustering Data With k -means and Single-Linkage Clustering

4.1 Clustering Dataset 1

When I cluster the ps1-clustering.txt dataset with both k -means and single-linkage clustering, there appear to be 7 clusters. This is confirmed by the figure [4.1](#) below. Single-linkage clustering perform better on this dataset, better and more easily separating the clusters, especially the center cluster. This occurs because k -means while only identify this as its own cluster if one of the cluster centers happened to be intialized to this center cluster. However, single-linkage clustering is not affected by this problem.

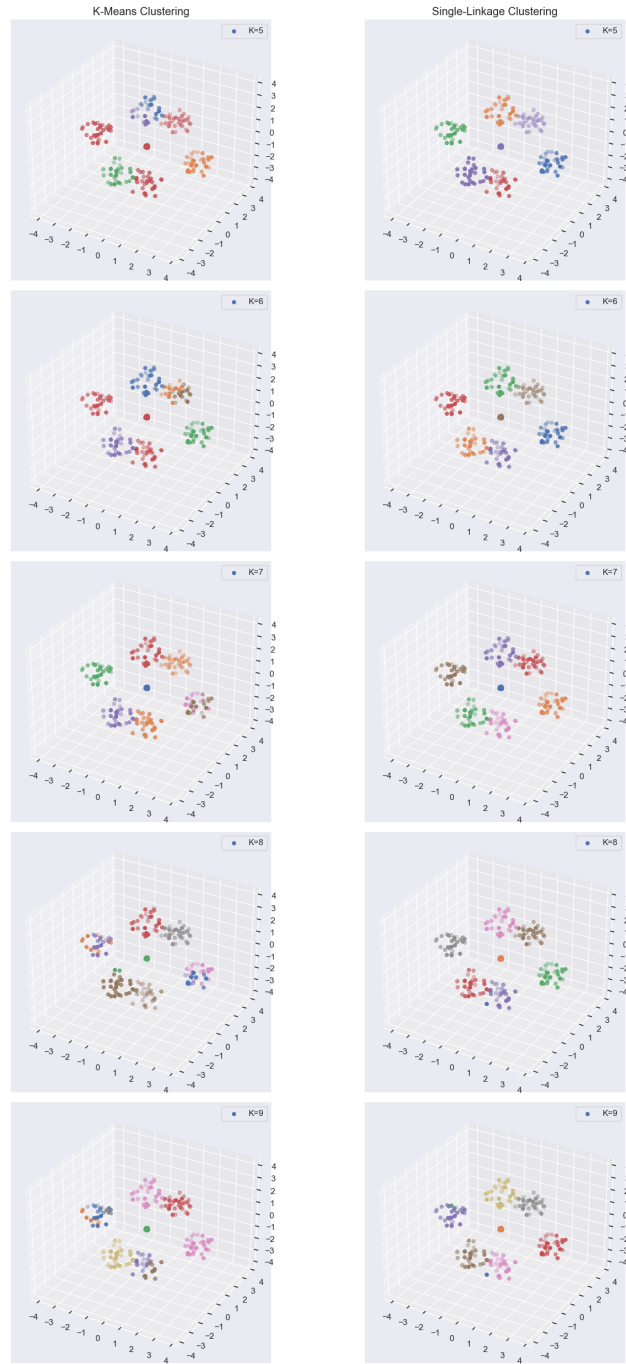


Figure 3: k -means and single-linkage clustering on ps1-clustering.txt

4.2 Clustering Dataset 2

This dataset appears to have no real structure to it – it is unclear how many clusters exist. However, single-linkage clustering often collapses into one cluster. k -means performs somewhat better, but it is not clear how many clusters to choose. This can be seen in the figure below [4.2](#) which clusters according to both k -means and single-linkage clustering for a variety of cluster values k to little success.

An interesting extension would be to measure the silhouette or Calinski-Harabasz score for each k and choose the k with the highest value score.

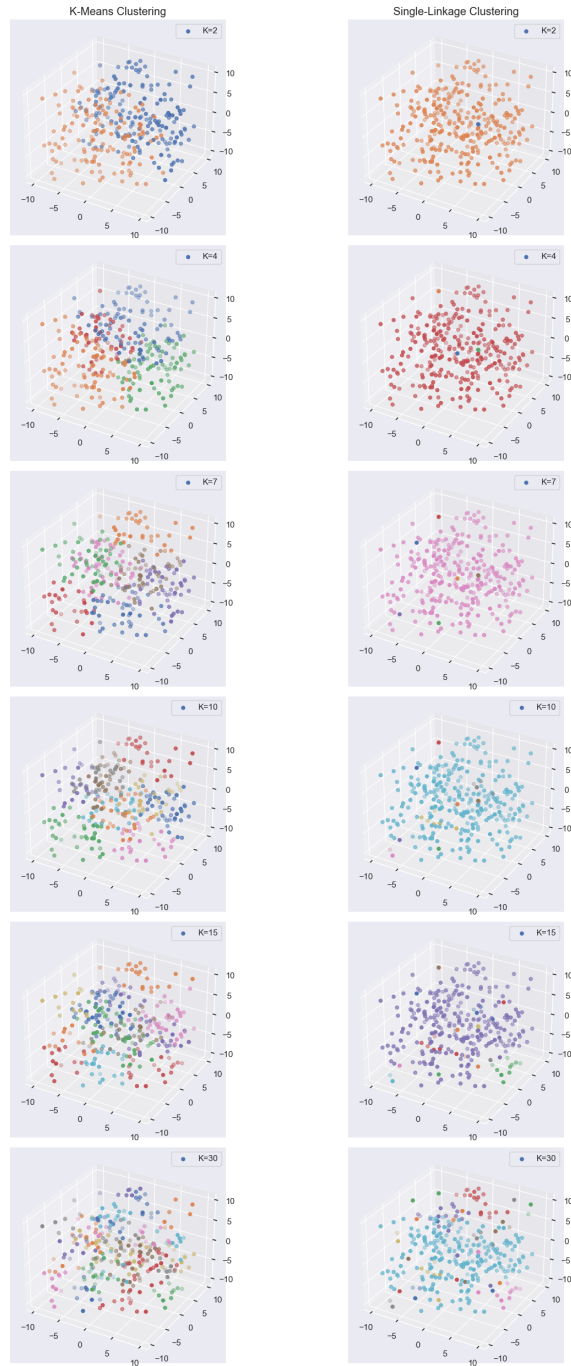


Figure 4: k -means and single-linkage clustering on ps1-data.txt

5 k -means Clustering Does Not Satisfy Consistency

This question invokes Jon Kleinberg’s famous 2002 paper, An Impossibility Theorem for Clustering. In it, he proves that no clustering algorithm can satisfy the following three properties: scale-invariance, richness, and consistency.

This question touches upon the consistency property, applied to the k -means algorithm: “suppose that the clustering Γ arises from the distance function d . If we now produce d' by reducing distances within the clusters and enlarging distance between the clusters then the same clustering Γ should arise from d' ”.

k -means trivially does not satisfy the consistency property, as do all centroid-based methods. For a simple example, let $k = 2$ and consider a set of points S . Divide S into two subsets, X with m data points and Y with $m\gamma$ data points for small $\gamma > 0$. Let the distance between points in X be r and the distance between points in Y be some small $\epsilon > 0$. Then let the distance between clusters X and Y be $r + \delta$ for some small $\delta > 0$. Then some initialization of k -means would select a point from X and Y which would result in the (true) clusters X and Y . However, consider splitting up X into sets X_0, X_1 with equal number of points. Then reduce the distance between points in X_0 to $r' < r$ and similarly reduce the distance between points in X_1 to $r' < r$. This changes the optimal choice of centroids to be points in X_0 and X_1 which would result in the (false) clusters X_0 and X_1 .

A similar type of intuition is captured in the figure below 5.




Figure 5: k -means unable to satisfy clustering consistency

6 Clustering Annuli With k -means and Spectral Clustering

Unsurprisingly, spectral clustering performs much better than k -means on the annuli dataset. This is because k -means uses Euclidean distance while spectral clustering creates a graph which captures local notions of distance. Thus, k -means will compute points in two different annuli as being much closer together than the spectral clustering algorithm. This can be seen in the figure below [6](#).




Figure 6: k -means vs spectral clustering on annuli data

7 Spectral Clustering: RBF Kernel vs K-Nearest Neighbors

Nearly done. The Gaussian method of constructing the weight matrix seems more robust.